



BASHAR ALMUNTASER

AI Engineer | Machine Learning | LLM Systems | Production AI | Backend Engineering

✉ almuntaserbashar@gmail.com ☎ +967 773771175 📍 Sana'a, Yemen
🔗 github.com/Bashar-1216 🔗 linkedin.com/in/bashar-almuntaser

Summary

AI Engineer specializing in the end-to-end architecture, deployment, and optimization of production-grade machine learning systems and LLM pipelines. Proven expertise in multi-model orchestration, bilingual (Arabic/English) NLP, and distributed microservices infrastructure. Adept at bridging complex AI research with scalable engineering, leveraging advanced techniques such as transformer fine-tuning, retrieval-augmented generation (RAG), and real-time data streaming to build resilient, high performance applications.

Education

2022 – 2026 **Bachelor of Science in Artificial Intelligence**
Sana'a, Yemen Emirates International University

Skills

Programming Python, TypeScript, SQL	AI & ML PyTorch, TensorFlow, Scikit-learn, Transformers, BERT, RoBERTa, Transfer Learning, LightGBM	LLM RAG, LLM APIs (OpenAI, Anthropic, Gemini, Perplexity), LangChain, AI Agents, Prompt Engineering, LLaMA-3, Ollama
NLP Arabic NLP, CAMEL-BERT, BiLSTM, NER, Sentiment Analysis, Text Classification, Domain-Adversarial Training	Data PySpark, Kafka, Pandas, NumPy, TimescaleDB, Elasticsearch	Backend FastAPI, NestJS, Celery, BullMQ, REST APIs
Frontend React, Next.js	DevOps Docker, Docker Compose, Linux, Git, Langfuse	Databases PostgreSQL, MongoDB, Redis

Languages

Arabic ● ● ● ● ● English ● ● ● ● ●

Projects

2025 – 2026 **GEO Platform**
Solo ML Engineer
Tech Stack: Python, NestJS, Next.js, PostgreSQL, Redis, Docker, OpenAI, Claude, Gemini, LangChain, Langfuse

- Architected an 8-stage asynchronous AI analysis pipeline executing structured prompt workflows across GPT-4, Claude, Gemini, and Perplexity using a Python worker architecture with BullMQ/Redis task queues and Pydantic schema validation.
- Engineered a bilingual Arabic/English entity resolution engine leveraging Unicode normalization, ALA-LC transliteration, legal suffix stripping, Jaro-Winkler fuzzy matching, and PostgreSQL trigram (pg_trgm) indexing, achieving 70–80% accuracy.

- Developed a hallucination detection and claim verification layer cross-referencing LLM-extracted brand claims against live web content, Wikidata SPARQL endpoints, and the Google Knowledge Graph API using multi signal confidence scoring.
- Designed a data-driven recommendation engine performing competitive feature gap analysis across Wikipedia, Schema.org, and Wikidata, scoring recommendations by Spearman rank correlation.
- Implemented production observability infrastructure including SHA-256 prompt version control, model drift detection via canary entity monitoring, and LLM cost/latency tracking with Langfuse and McNemar's test.

2025 – 2026

Smart Amazon Product Analyzer

Solo Full-Stack ML Engineer

Tech Stack: Python, FastAPI, React, TypeScript, TimescaleDB, Elasticsearch, Redis, Celery, Docker, LightGBM, BERT, LLaMA-3, Ollama

- Engineered a five-indicator product scoring engine combining LightGBM demand forecasting with a hybrid BERT and LLaMA-3 NLP pipeline (via Ollama) for review toxicity detection and competitive market analysis (HHI).
- Designed and implemented a Margin Kill-Switch automation rejecting investment opportunities when net ROI falls below thresholds based on real-time Amazon FBA fee calculations.
- Deployed a resilient 8-container Docker Compose microservices architecture on Linux servers with health check-ordered startup sequencing, ensuring high availability and scalable backend services.
- Built an executive analytics dashboard with radar charts and time-series visualizations classifying products as RECOMMENDED, MONITOR, or REJECT based on institutional investment thresholds.

2024 – 2024

Financial Fraud Detection Platform

ML Engineer

Tech Stack: Python, PySpark, Apache Kafka, MongoDB, Scikit-learn

- Built an end-to-end fraud detection pipeline using PySpark and ensemble machine learning models to classify anomalous financial transactions at scale, handling structured and unstructured data.
- Engineered fraud-signal features and simulated real-time transaction streaming using Apache Kafka, storing model predictions in MongoDB for downstream tracking and analysis.

2024 – 2024

Sentiment Analysis & Fake Review Detection

ML Engineer

Tech Stack: Python, CAMEL-BERT, BiLSTM, Transformers, FastAPI, Domain-Adversarial Training

- Fine-tuned CAMEL-BERT and BiLSTM architectures for domain-specific Arabic sentiment classification, building custom Arabic text preprocessing pipelines to handle social media dialects and cross-domain generalization.
- Developed a domain-adversarial transformer-based NLP system for cross-category fake review detection, deployed as a real-time inference service using FastAPI.